

**Problem statement:**

Visualise how data is distributed as RDD across a cluster of 5 machines of 4 cores. Describe a complete picture of RDD, partitions, lambda implementing on each partition and correlation with the JVMs.

**Walkthrough:**

RDD's are:

Resilient - Fault tolerant (immutable)

Distributed – Divided across the clusters

Datasets – collection of records.

The Spark automatically partitions RDDs and distribute the partitions across different nodes. The RDDs and its partitions define the basic units of parallelism in spark.

In local mode the default no. of partitions is the number of cores allotted to the application.

While in cluster mode, the no. of partitions depends on no. of blocks or input splits.

So by default parallelism in spark is no. of cores allocated in sc.

**Snapshots:**

I have extracted the image behind executors and partitions in relation with JVM, the options to check micropartitions and internal functioning through databricks is not proper , do have a look at snapshots and guide me.

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

life lessons through... Online C Compiler ... general | Xebia-UPe... C program to insert... turin6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

Clusters / sparkCluster

sparkCluster Edit Clone Restart Terminate Delete

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

stage 5

Scan csv

FileSourceRDD [25]  
rdd at command-25964311794340491

MapPartitionsRDD [26]  
rdd at command-25964311794340491

DeserializeToObject

MapPartitionsRDD [27]  
rdd at command-25964311794340491

SQLDecoderRDD [28]  
rdd at command-25964311794340491

mapPartitions

MapPartitionsRDD [29]  
rdd at command-25964311794340491

glom

MapPartitionsRDD [30]  
glom at command-25964311794340501

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

life lessons through... Online C Compiler ... general | Xebia-UPe... C program to insert... turin6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

Clusters / sparkCluster

sparkCluster Edit Clone Restart Terminate Delete

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

☒ Result Serialization Time  
☒ Getting Result Time  
☒ Peak Execution Memory  
☒ Peak JVM Memory OnHeap / OffHeap  
☒ Peak Execution Memory OnHeap / OffHeap  
☒ Peak Storage Memory OnHeap / OffHeap  
☒ Peak Pool Memory Direct / Mapped

Event Timeline

Summary Metrics for 8 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Task Deserialization Time	0.1 s	0.1 s	0.2 s	0.2 s	0.2 s
Duration	0.4 s	0.4 s	0.6 s	0.6 s	0.6 s
GC Time	0.0 ms	0.0 ms	0.0 ms	0.0 ms	0.0 ms
Result Serialization Time	0.0 ms	0.0 ms	0.0 ms	2.0 ms	4.0 ms
Getting Result Time	0.0 ms	0.0 ms	0.0 ms	0.0 ms	0.0 ms
Scheduler Delay	46.0 ms	48.0 ms	80.0 ms	0.1 s	0.1 s
Peak Execution Memory	0.0 B	0.0 B	0.0 B	0.0 B	0.0 B

Showing 1 to 7 of 7 entries

Aggregated Metrics by Executor

Show 20 entries

Executor ID	Logs	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks	Excluded	Peak JVM Memory OnHeap / OffHeap	Peak Execution Memory OnHeap / OffHeap	Peak Storage Memory OnHeap / OffHeap	Peak Pool Memory Direct / Mapped
driver		10.172.242.170:4473	6 s	8	0	0	8	false	1.7 GiB / 311.8 MiB	0.0 B / 0.0 B	2.8 MiB / 0.0 B	3.6 MiB / 0.0 B

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

driver 10.172.242.170:44073 6 s 8 0 0 8 false 1.7 GiB / 311.8 MiB 0.0 B / 0.0 B 2.8 MiB / 0.0 B 3.6 MB / 0.0 B

Showing 1 to 1 of 1 entries

Tasks (8)

Show 20 entries

Index	Task ID	Attempt	Status	Locality level	Executor ID	Host	Logs	Launch Time	Duration	GC Time	Scheduler Delay	Task Deserialization Time	Result Serialization Time	Getting Result Time	Peak Execution Memory	Errors
0	17	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.4 s		48.0 ms	0.2 s				
1	18	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.6 s		46.0 ms	0.2 s				
2	19	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.6 s		46.0 ms	0.2 s				
3	20	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.6 s		77.0 ms	0.2 s				
4	21	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.4 s		80.0 ms	0.2 s	4.0 ms			
5	22	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.4 s		0.1 s	0.1 s				
6	23	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.5 s		0.1 s	0.1 s				
7	24	0	SUCCESS	PROCESS_LOCAL	driver	ip-10-172-242-170.us-west-2.compute.internal		2021-05-18 12:58:16	0.6 s		0.1 s	0.1 s	2.0 ms			

Showing 1 to 8 of 8 entries

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

Event Timeline

DAAG Visualization

```
graph TD; A[Scan csv] --> B[DeserializeToObject]; B --> C[mapPartitions]; C --> D[glom];
```

Amazon WorkMail x RN11 board - Agile board x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

### Executors

Show Additional Metrics

Summary

	RDD Blocks	Storage Memory	On Heap Storage Memory	Off Heap Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	218.1 KiB / 3.9 GiB	218.1 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	25	25	52 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	218.1 KiB / 3.9 GiB	218.1 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	25	25	52 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0

### Executors

Show 20 entries

Executor ID	Address	Status	RDD Blocks	Storage Memory	On Heap Storage Memory	Off Heap Storage Memory	Peak JVM Memory OnHeap / OffHeap	Peak Execution Memory OnHeap / OffHeap	Peak Storage Memory OnHeap / OffHeap	Peak Pool Memory Direct / Mapped	Disk Used	Cores	Resources	Resource Profile Id	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump	Heap Histogram
driver	10.172.242.170:44073	Active	0	218.1 KiB / 3.9 GiB	218.1 KiB / 3.9 GiB	0.0 B / 0.0 B	2.3 GiB / 312.9 MiB	0.0 B / 0.0 B	2.8 MiB / 0.0 B	4.6 MiB / 0.0 B	0.0 B	8			0	0	25	25	52 s (0.0 ms)	0.0 B	0.0 B	0.0 B	Thread Dump	Heap Histogram

Showing 1 to 1 of 1 entries

Amazon WorkMail x RN11 board - Agile board x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

### Spark Jobs (7)

User: root

Total Uptime: 2.2 h

Scheduling Mode: FAIR

Completed Jobs: 6

Event Timeline

Completed Jobs (6)

Page: 1

1 Pages. Jump to 1. Show 100 items in a page. Go

Job Id (Job Group)	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5 (263617408804822260,5383678818898984307,79747d9f196429c8676b3f5a9a8443a)	a.rddglom().collect()	2021/05/18 07:28:16	0.9 s	1/1	8/8
4 (263617408804822260,6485781451206622927,832b22446ee44a2191cd8690e010762b)	ashow()	2021/05/18 07:28:05	0.4 s	1/1	3/3
3 (263617408804822260,6485781451206622927,832b22446ee44a2191cd8690e010762b)	ashow()	2021/05/18 07:28:05	0.5 s	1/1	4/4
2 (263617408804822260,6485781451206622927,832b22446ee44a2191cd8690e010762b)	ashow()	2021/05/18 07:28:05	0.4 s	1/1	1/1
1 (263617408804822260,778165427652549013,173ac723169c4fdb36af3a7bea5c2ea)	val a = spark.read.csv("/Users/lakshay sharma/D... csv at command-2986431179434048:1	2021/05/18 07:24:12	0.8 s	1/1	1/1
0 (263617408804822260,5445176125542806598,691191ac94984cdeb90aa035fb49312d)	numberDF.write.csv("/Users/lakshay sharma/Docum... csv at command-2986431179434048:1	2021/05/18 07:03:29	6 s	1/1	8/8

Page: 1

1 Pages. Jump to 1. Show 100 items in a page. Go

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps life lessons through... Online C Compiler... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

### Clusters / sparkCluster

sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs	Stages	Storage	Environment	Executors	SQL	JDBC/ODBC Server	Structured Streaming								
Active(1)	0	89.4 KiB / 3.9 GiB	89.4 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	9	9	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	89.4 KiB / 3.9 GiB	89.4 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	9	9	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0

#### Executors

Show 20 entries Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	On Heap Storage Memory	Off Heap Storage Memory	Peak JVM Memory OnHeap / OffHeap	Peak Execution Memory OnHeap / OffHeap	Peak Storage Memory OnHeap / OffHeap	Peak Pool Memory Direct / Mapped	Disk Used	Cores	Resources	Resource Profile Id	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks
driver	10.172.242.170:44073	Active	0	89.4 KiB / 3.9 GiB	89.4 KiB / 3.9 GiB	0.0 B / 0.0 B	2.3 GiB / 305.3 MiB	0.0 B / 0.0 B	1.2 MiB / 0.0 B	3.1 MiB / 0.0 B	0.0 B	8			0	0	9	9

Showing 1 to 1 of 1 entries Previous Next

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps life lessons through... Online C Compiler... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

### Clusters / sparkCluster

sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

#### Details for Stage 5 (Attempt 0)

Resource Profile Id: 0  
Total Time Across All Tasks: 4 s  
Locality Level Summary: Process local: 8  
Associated Job Ids: 5

DAG Visualization  
Show Additional Metrics  
Event Timeline  
Enable zooming

Scheduler Delay Task Deserialization Time Shuffle Write Time Shuffle Read Time Executor Computing Time Shuffle Write Time Result Serialization Time Getting Result Time

Tasks: 8. 1 Pages. Jump to 1. Show 8 items in a page. Go

Summary Metrics for 8 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
...	...	...	...	...	...

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps Life lessons through... Online C Compiler ... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

### Clusters / sparkCluster

sparkCluster [Edit] [Clone] [Restart] [Terminate] [Delete]

Configuration Notebooks (1) Libraries Event Log **Spark UI** Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment **Executors** SQL JDBC/ODBC Server Structured Streaming

#### Executors

Show Additional Metrics

##### Summary

	RDD Blocks	Storage Memory	On Heap Storage Memory	Off Heap Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	89.4 KiB / 3.9 GiB	89.4 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	9	9	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	89.4 KiB / 3.9 GiB	89.4 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	9	9	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0

#### Executors

Show 20 entries Search:

	On Heap	Off Heap	Peak JVM Memory OnHeap	Peak Execution Memory	Peak Storage Memory OnHeap	Peak Pool Memory
--	---------	----------	------------------------	-----------------------	----------------------------	------------------

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps Life lessons through... Online C Compiler ... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

### Clusters / sparkCluster

sparkCluster [Edit] [Clone] [Restart] [Terminate] [Delete]

Configuration Notebooks (1) Libraries Event Log **Spark UI** Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors **SQL** JDBC/ODBC Server Structured Streaming

#### Spark Jobs (?)

User: root  
Total Uptime: 2.2 h  
Scheduling Mode: FAIR  
Completed Jobs: 2

Event Timeline

##### Completed Jobs (2)

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id (Job Group)	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1 (263617408804822260_7781654276525249013_173ac723169c4fdbb36af3a7bea5c2ea)	val a = spark.read.csv("/Users/lakshay sharma/D... csv at command-2986431179434047:1	2021/05/18 07:24:12	0.8 s	1/1	1/1
0 (263617408804822260_5445176125542806598_691191ac94984cdeb90aa035fb49312d)	numberDF.write.csv("/Users/lakshay sharma/Docum... csv at command-2986431179434046:1	2021/05/18 07:03:29	6 s	1/1	8/8

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Amazon WorkMail x RNIT board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps life lessons through... Online C Compiler ... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

88	DBFSV1-HttpClient-88	RUNNABLE	Monitor(sun.nio.ch.Util\$3@1864542920), Monitor(java.util.Collections\$UnmodifiableSet@2018685279), Monitor(sun.nio.ch.EPollSelectorImpl@281836090)
89	DBFSV1-HttpClient-89	TIMED_WAITING	
90	DBFSV1-HttpClient-90	TIMED_WAITING	
124	DestroyJavaVM	RUNNABLE	
47	dispatcher-BlockManagerEndpoint1	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@571696027))
46	dispatcher-BlockManagerMaster	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1053719038))
30	dispatcher-event-loop-0	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1287215032))
31	dispatcher-event-loop-1	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1021082377))
32	dispatcher-event-loop-2	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1935994577))
33	dispatcher-event-loop-3	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1749978269))
34	dispatcher-event-loop-4	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1858863803))
35	dispatcher-event-loop-5	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@831993944))
36	dispatcher-event-loop-6	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1699826397))
37	dispatcher-event-loop-7	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@2045188238))

Windows taskbar: 12:52 PM 18-05-2021

Amazon WorkMail x RNIT board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps life lessons through... Online C Compiler ... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log Spark UI Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL JDBC/ODBC Server Structured Streaming

### Thread dump for executor driver

Updated at 2021/05/18 07:22:16

Expand All

Search:

Thread ID	Thread Name	Thread State	Thread Locks
428	Attach Listener	RUNNABLE	
170	BoneCP-keep-alive-scheduler	TIMED_WAITING	
173	BoneCP-keep-alive-scheduler	TIMED_WAITING	
171	BoneCP-pool-watch-thread	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@186548010))
174	BoneCP-pool-watch-thread	WAITING	Lock(java.util.concurrent.ThreadPoolExecutor\$Worker@1517072245))
169	com.google.common.base.internal.Finalizer	WAITING	
172	com.google.common.base.internal.Finalizer	WAITING	
126	Connector-Scheduler-25a2c4dc-1	TIMED_WAITING	
71	context-cleaner-periodic-gc	TIMED_WAITING	

Windows taskbar: 12:52 PM 18-05-2021

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps life lessons through... Online C Compiler ... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

### Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log **Spark UI** Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment **Executors** SQL JDBC/ODBC Server Structured Streaming

Active(1)	0	22.3 KiB / 3.9 GiB	22.3 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	8	8	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	22.3 KiB / 3.9 GiB	22.3 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	8	8	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0

#### Executors

Show 20 entries Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	On Heap Storage Memory	Off Heap Storage Memory	Peak JVM Memory OnHeap / OffHeap	Peak Execution Memory OnHeap / OffHeap	Peak Storage Memory OnHeap / OffHeap	Peak Pool Memory Direct / Mapped	Disk Used	Cores	Resources	Resource Profile Id	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks
driver	10.172.242.170:44073	Active	0	22.3 KiB / 3.9 GiB	22.3 KiB / 3.9 GiB	0.0 B / 0.0 B	2.3 GiB / 295.7 MiB	0.0 B / 0.0 B	807.3 KiB / 0.0 B	1.4 MiB / 0.0 B	0.0 B	8		0	0	0	8	8

Showing 1 to 1 of 1 entries Previous Next

Amazon WorkMail x RN11 board - Agile board - x (155) How Spark Creates P... x task1 - Databricks Commu... x what does thread dump m... x Clusters - Databricks Comm... x

community.cloud.databricks.com/?o=1681350184697554#setting/clusters/0518-051549-sunk595/sparkUI

Apps life lessons through... Online C Compiler ... general | Xebia-UPE... C program to insert... turing6 What Defines an Eff... Stay Disciplined C program for bina... CBSE Class 11 Che... Reading list

### Clusters / sparkCluster

Configuration Notebooks (1) Libraries Event Log **Spark UI** Driver Logs Metrics Apps Spark Cluster UI - Master

Hostname: ec2-54-185-87-23.us-west-2.compute.amazonaws.com Spark Version: 8.2.x-scala2.12

Jobs Stages Storage Environment **Executors** SQL JDBC/ODBC Server Structured Streaming

#### Executors

Show Additional Metrics

##### Summary

	RDD Blocks	Storage Memory	On Heap Storage Memory	Off Heap Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	22.3 KiB / 3.9 GiB	22.3 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	8	8	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	22.3 KiB / 3.9 GiB	22.3 KiB / 3.9 GiB	0.0 B / 0.0 B	0.0 B	8	0	0	8	8	43 s (0.0 ms)	0.0 B	0.0 B	0.0 B	0

#### Executors

Show 20 entries Search:

	On Heap	Off Heap	Peak JVM Memory OnHeap	Peak Execution Memory	Peak Storage Memory OnHeap	Peak Pool Memory
--	---------	----------	------------------------	-----------------------	----------------------------	------------------



